

## CLONES

A person can interact with your service and be served at any request of any server – application servers involve a lot of servers.

– Every server contains exactly the same codebase and does not store any user-related data, like sessions or profile pictures.

– Sessions need to be stored in a centralized data store – accessible to all application servers:

- + Either external database
- + Or external persistent cache (Redis) – better performance than external database

External = data store does not reside on the application servers – somewhere in or near the data center of your application servers

### – Why This Matters?

- **Session Consistency:** Users can be routed to any application server, and their session state will remain consistent.
- **Load Balancing:** Session data is not tied to a specific server, allowing load balancers to distribute requests efficiently.
- **Fault Tolerance:** If an application server goes down, session data is not lost, as it is stored centrally.

### – Use Case Scenarios

- **High-Traffic Applications:** Websites with many users, where user sessions need to be maintained across various servers.
- **Microservices Architectures:** Different services need to share session data to provide a seamless user experience.
- **Cloud Environments:** Applications deployed in cloud environments often use centralized data stores for session management due to their distributed nature.

– How do we make sure that a code change is sent to all of the servers without one server still serving old code? – Capistrano: a deployment automation tool primarily designed for Ruby on Rails applications but can be adapted for other technologies.

## **Steps After Outsourcing Sessions and Standardizing the Codebase:**

### **Outsource Sessions:**

- Session data is stored in a centralized external data store (e.g., Redis or a database).
- This decouples session management from the application servers, allowing them to be stateless.

### **Serve the Same Codebase:**

- Ensure all servers run the same code version by automating deployment with Capistrano.

### **Create an Amazon Machine Image (AMI):**

– An Amazon Machine Image is a snapshot of an EC2 instance that includes the operating system, application code, runtime, libraries, and configurations.

#### **– Steps to Create an AMI:**

- + Deploy your application on a server.
- + Ensure it is fully configured and running the latest code.
- + Create an AMI from this server, capturing its state.

#### **– Use the AMI as a Base for New Instances:**

- + New EC2 instances can be launched using the AMI, ensuring they start with the same base setup.
- + This process guarantees consistency in the environment and application state across all instances.

#### **– Initial Deployment of Latest Code:**

- + After launching a new instance from the AMI, **perform an initial deployment using Capistrano to ensure the instance has the latest code changes.**
- + This step is crucial as the AMI might not always include the very latest code if updates were made after the image was created.